



Problem Statement:

Prepare a classification model using Naive Bayes for salary data test data

1. Import Necessary Libraries

```
In [1]: import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
```

2. Import Data

```
In [3]: Naive_bayes_Test_Data = pd.read_csv('SalaryData_Test.csv')
Naive_bayes_Test_Data
```

Out[3]:

	age	workclass	education	educationno	maritalstatus	occupation	relationship	race
0	25	Private	11th	7	Never-married	Machine-op-inspct	Own-child	Black
1	38	Private	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White
2	28	Local-gov	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	White
3	44	Private	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Black
4	34	Private	10th	6	Never-married	Other-service	Not-in-family	White
...
15055	33	Private	Bachelors	13	Never-married	Prof-specialty	Own-child	White
15056	39	Private	Bachelors	13	Divorced	Prof-specialty	Not-in-family	White
15057	38	Private	Bachelors	13	Married-civ-spouse	Prof-specialty	Husband	White
15058	44	Private	Bachelors	13	Divorced	Adm-clerical	Own-child	Asian-Pac-Islander
15059	35	Self-emp-inc	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White

15060 rows × 14 columns





3. Data Understanding

```
In [5]: Naive_bayes_Test_Data.shape
```

```
Out[5]: (15060, 14)
```

```
In [6]: Naive_bayes_Test_Data.dtypes
```

```
Out[6]: age                int64
workclass                object
education                object
educationno              int64
maritalstatus            object
occupation                object
relationship              object
race                    object
sex                     object
capitalgain              int64
capitalloss              int64
hoursperweek             int64
native                   object
Salary                   object
dtype: object
```

```
In [7]: Naive_bayes_Test_Data.isna().sum()
```

```
Out[7]: age                0
workclass                0
education                0
educationno              0
maritalstatus            0
occupation                0
relationship              0
race                    0
sex                     0
capitalgain              0
capitalloss              0
hoursperweek             0
native                   0
Salary                   0
dtype: int64
```

In [8]: Naive_bayes_Test_Data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 15060 entries, 0 to 15059
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   age                    15060 non-null  int64
1   workclass              15060 non-null  object
2   education              15060 non-null  object
3   educationno            15060 non-null  int64
4   maritalstatus          15060 non-null  object
5   occupation             15060 non-null  object
6   relationship           15060 non-null  object
7   race                   15060 non-null  object
8   sex                    15060 non-null  object
9   capitalgain            15060 non-null  int64
10  capitalloss            15060 non-null  int64
11  hoursperweek           15060 non-null  int64
12  native                  15060 non-null  object
13  Salary                  15060 non-null  object
dtypes: int64(5), object(9)
memory usage: 1.6+ MB
```

In [9]: Naive_bayes_Test_Data.describe(include='all')

Out[9]:

	age	workclass	education	educationno	maritalstatus	occupation	relationship
count	15060.000000	15060	15060	15060.000000	15060	15060	15060
unique	NaN	7	16	NaN	7	14	6
top	NaN	Private	HS-grad	NaN	Married-civ-spouse	Exec-managerial	Husband
freq	NaN	11021	4943	NaN	6990	1992	6203
mean	38.768327	NaN	NaN	10.112749	NaN	NaN	NaN
std	13.380676	NaN	NaN	2.558727	NaN	NaN	NaN
min	17.000000	NaN	NaN	1.000000	NaN	NaN	NaN
25%	28.000000	NaN	NaN	9.000000	NaN	NaN	NaN
50%	37.000000	NaN	NaN	10.000000	NaN	NaN	NaN
75%	48.000000	NaN	NaN	13.000000	NaN	NaN	NaN
max	90.000000	NaN	NaN	16.000000	NaN	NaN	NaN

```
In [10]: Naive_bayes_Test_Data.describe(include='all').nunique()
```

```
Out[10]: age                8
workclass            4
education            4
educationno          8
maritalstatus        4
occupation            4
relationship          4
race                 4
sex                  4
capitalgain           5
capitalloss           5
hoursperweek          7
native                4
Salary                4
dtype: int64
```

```
In [12]: Naive_bayes_Test_Data['workclass'].nunique()
```

```
Out[12]: 7
```

```
In [11]: Naive_bayes_Test_Data['workclass'].unique()
```

```
Out[11]: array([' Private', ' Local-gov', ' Self-emp-not-inc', ' Federal-gov',
                ' State-gov', ' Self-emp-inc', ' Without-pay'], dtype=object)
```

In [13]: Naive_bayes_Test_Data

Out[13]:

	age	workclass	education	educationno	maritalstatus	occupation	relationship	race
0	25	Private	11th	7	Never-married	Machine-op-inspct	Own-child	Black
1	38	Private	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White
2	28	Local-gov	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	White
3	44	Private	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Black
4	34	Private	10th	6	Never-married	Other-service	Not-in-family	White
...
15055	33	Private	Bachelors	13	Never-married	Prof-specialty	Own-child	White
15056	39	Private	Bachelors	13	Divorced	Prof-specialty	Not-in-family	White
15057	38	Private	Bachelors	13	Married-civ-spouse	Prof-specialty	Husband	White
15058	44	Private	Bachelors	13	Divorced	Adm-clerical	Own-child	Asian-Pac-Islander
15059	35	Self-emp-inc	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White

15060 rows × 14 columns



4. Data Preparation

In [16]: `from sklearn.preprocessing import LabelEncoder`

In [17]: `Label = LabelEncoder()`

In [18]: `Naive_bayes_Test_Data['workclass'] = Label.fit_transform(Naive_bayes_Test_Data['workclass'])`
`Naive_bayes_Test_Data['education'] = Label.fit_transform(Naive_bayes_Test_Data['education'])`
`Naive_bayes_Test_Data['maritalstatus'] = Label.fit_transform(Naive_bayes_Test_Data['maritalstatus'])`
`Naive_bayes_Test_Data['occupation'] = Label.fit_transform(Naive_bayes_Test_Data['occupation'])`
`Naive_bayes_Test_Data['relationship'] = Label.fit_transform(Naive_bayes_Test_Data['relationship'])`
`Naive_bayes_Test_Data['race'] = Label.fit_transform(Naive_bayes_Test_Data['race'])`
`Naive_bayes_Test_Data['sex'] = Label.fit_transform(Naive_bayes_Test_Data['sex'])`
`Naive_bayes_Test_Data['native'] = Label.fit_transform(Naive_bayes_Test_Data['native'])`
`Naive_bayes_Test_Data['Salary'] = Label.fit_transform(Naive_bayes_Test_Data['Salary'])`



In [19]: Naive_bayes_Test_Data

Out[19]:

	age	workclass	education	educationno	maritalstatus	occupation	relationship	race	sex
0	25	2	1	7	4	6	3	2	1
1	38	2	11	9	2	4	0	4	1
2	28	1	7	12	2	10	0	4	1
3	44	2	15	10	2	6	0	2	1
4	34	2	0	6	4	7	1	4	1
...
15055	33	2	9	13	4	9	3	4	1
15056	39	2	9	13	0	9	1	4	0
15057	38	2	9	13	2	9	0	4	1
15058	44	2	9	13	0	0	3	1	1
15059	35	3	9	13	2	3	0	4	1

15060 rows × 14 columns



In [20]: Naive_bayes_Test_Data.dtypes

Out[20]:

age	int64
workclass	int32
education	int32
educationno	int64
maritalstatus	int32
occupation	int32
relationship	int32
race	int32
sex	int32
capitalgain	int64
capitalloss	int64
hoursperweek	int64
native	int32
Salary	int32
dtype:	object

5. Model Building

In [21]:

```
X = Naive_bayes_Test_Data.drop(['Salary'], axis=1)
y = Naive_bayes_Test_Data['Salary']
```

In [22]: X

Out[22]:

	age	workclass	education	educationno	maritalstatus	occupation	relationship	race	sex
0	25	2	1	7	4	6	3	2	1
1	38	2	11	9	2	4	0	4	1
2	28	1	7	12	2	10	0	4	1
3	44	2	15	10	2	6	0	2	1
4	34	2	0	6	4	7	1	4	1
...
15055	33	2	9	13	4	9	3	4	1
15056	39	2	9	13	0	9	1	4	0
15057	38	2	9	13	2	9	0	4	1
15058	44	2	9	13	0	0	3	1	1
15059	35	3	9	13	2	3	0	4	1

15060 rows × 13 columns

In [23]: y

```
Out[23]: 0      0
1      0
2      1
3      1
4      0
..
15055   0
15056   0
15057   0
15058   0
15059   1
Name: Salary, Length: 15060, dtype: int32
```

In [47]: `from sklearn.model_selection import train_test_split`In [49]: `X_train, X_test, y_train, y_test = train_test_split(X,y, train_size=0.20, random_`In [50]: `X_train.shape, y_train.shape`Out[50]: `((3012, 13), (3012,))`In [51]: `X_test.shape, y_test.shape`Out[51]: `((12048, 13), (12048,))`



6. Model training | Testing | Evaluation

```
In [52]: from sklearn.naive_bayes import MultinomialNB
nb_Classifier = MultinomialNB()
```

```
In [59]: nb_Classifier.fit(X_test, y_test)
MultinomialNB()
```

Out[59]: MultinomialNB()

```
In [60]: y_pred = nb_Classifier.predict(X_train)
y_pred
```

Out[60]: array([0, 0, 0, ..., 0, 0, 0])

```
In [61]: from sklearn.metrics import accuracy_score, precision_score, confusion_matrix, cla
```

```
In [63]: print("Accuracy Score:", round(accuracy_score(y_train, y_pred),4))
```

Accuracy Score: 0.7915

```
In [64]: print("Confusion Matrix:", confusion_matrix(y_train, y_pred))
```

Confusion Matrix: [[2241 75]
[553 143]]

```
In [65]: print("Classification Report:", classification_report(y_train, y_pred))
```

Classification Report:			precision	recall	f1-score	support
0	0.80	0.97	0.88		2316	
1	0.66	0.21	0.31		696	
accuracy			0.79		3012	
macro avg			0.73	0.59	0.60	3012
weighted avg			0.77	0.79	0.75	3012

```
In [66]: print("Precision Score:", round(precision_score(y_train, y_pred),4))
```

Precision Score: 0.656

7. Model Deployment

```
In [73]: from pickle import dump
```

```
In [74]: dump(nb_Classifier, open('lognb_Classifier.pkltest', 'wb'))
```

```
In [75]: from pickle import load
```




```
In [76]: nb_Classifier_pickle = load(open('lognb_Classifier.plktest', 'rb'))
```

```
In [77]: Pickle_pred = nb_Classifier_pickle.predict(X_train)
```

Conclusion:

Classification model using Naive Bayes for salary data train data with model accuracy 79%